

INTRO TO DATA SCIENCE

TOPICS IN DATA ENGINEERING

INTRO TO DATA SCIENCE

BIG DATA

BIG DATA

What does *big data* actually refer to?

BIG DATA: FUN FACTS

The Large Hadron Collider (LHC) produces roughly 30 petabytes per year

The *LHC Data Centre*:

- has 11,000 servers with 100,000 processor cores
- processes about one petabyte of data every day - the equivalent of around 210,000 DVDs

source: [CERN Computing](#)

BIG DATA

What does ***big data*** actually refer to?

A common approach is to talk about big data in terms of the 3 Vs:

- Volume
- Velocity
- Variety

source: [3-D Data Management: Controlling Data Volume, Velocity and Variety](#) (Feb 6, 2001)

BIG DATA

- ***Volume***

- more than can fit in memory on your laptop, e.g. Amazon user behavior data

- ***Velocity***

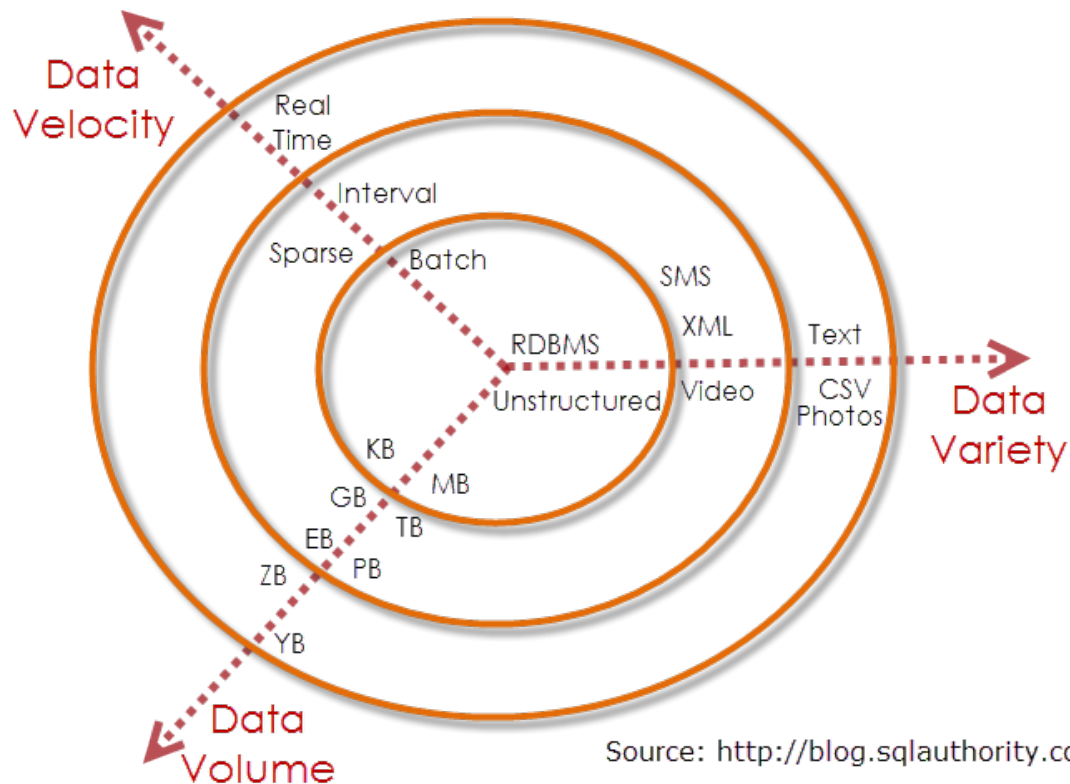
- faster than standard machines can process, e.g. Twitter Firehose

- ***Variety***

- does not conform to a single structure, e.g. Google's cache of web pages

BIG DATA

3Vs of Big Data



Source: <http://blog.sqlauthority.com>

SCALING BIG DATA PROCESSING

How would you approach dealing with this kind of data?

One approach would be to get a huge supercomputer, but this has some obvious drawbacks:

- expensive
- difficult to maintain
- scalability is bounded

SCALING BIG DATA PROCESSING

How would you approach dealing with this kind of data?

One approach would be to get a huge supercomputer, but this has some obvious drawbacks:

- expensive
- difficult to maintain
- ***scalability is bounded***

SCALING BIG DATA PROCESSING

Instead of one huge machine, what if we got a bunch of regular (commodity) machines?

- cheaper
- easier to maintain
- scalability is unbounded
 - “just add more nodes to the cluster”

SCALING BIG DATA PROCESSING

Guiding Principles:

- Scale horizontally, not vertically
- Assume hardware failure is common
- Move code to the data
- Hide system level-details from developers (and data scientists)
- Seamless scalability

INTRO TO DATA SCIENCE

MAPREDUCE

SCALING BIG DATA PROCESSING

One day when I was having lunch with Richard Feynman, I mentioned to him that I was planning to start a company to build a parallel computer with a million processors. His reaction was unequivocal,

"That is positively the dopiast idea I ever heard."

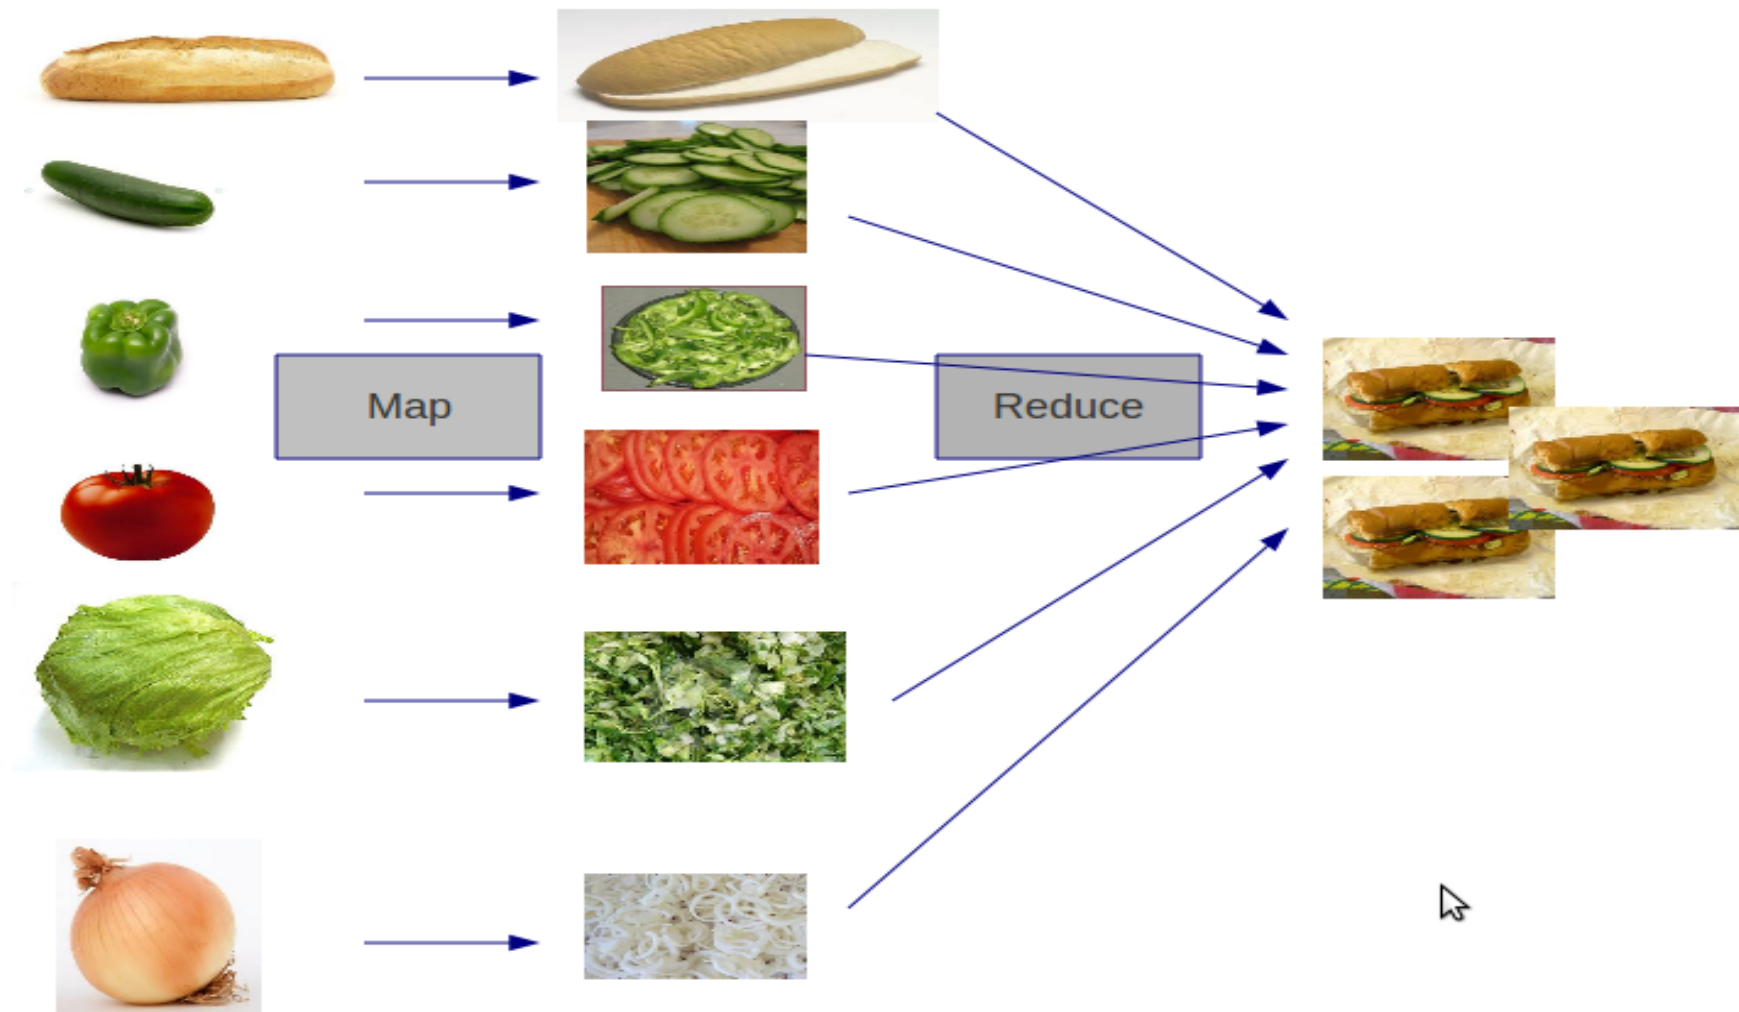
- Danny Hillis in *Physics Today* on the Connection Machine,
reposted: [Richard Feynman and The Connection Machine](#)

MAPREDUCE

From the famous [Google Research Paper](#):

MapReduce: Simplified Data Processing on Large Clusters

- A programming model and an associated implementation for processing and generating large data sets
- MapReduce programs are automatically parallelized and executed on a large cluster of commodity machines



MAPREDUCE

MapReduce involves splitting a problem into subtasks and processing these subtasks in parallel.

This takes place in two phases:

1. the mapper phase
2. the reducer phase

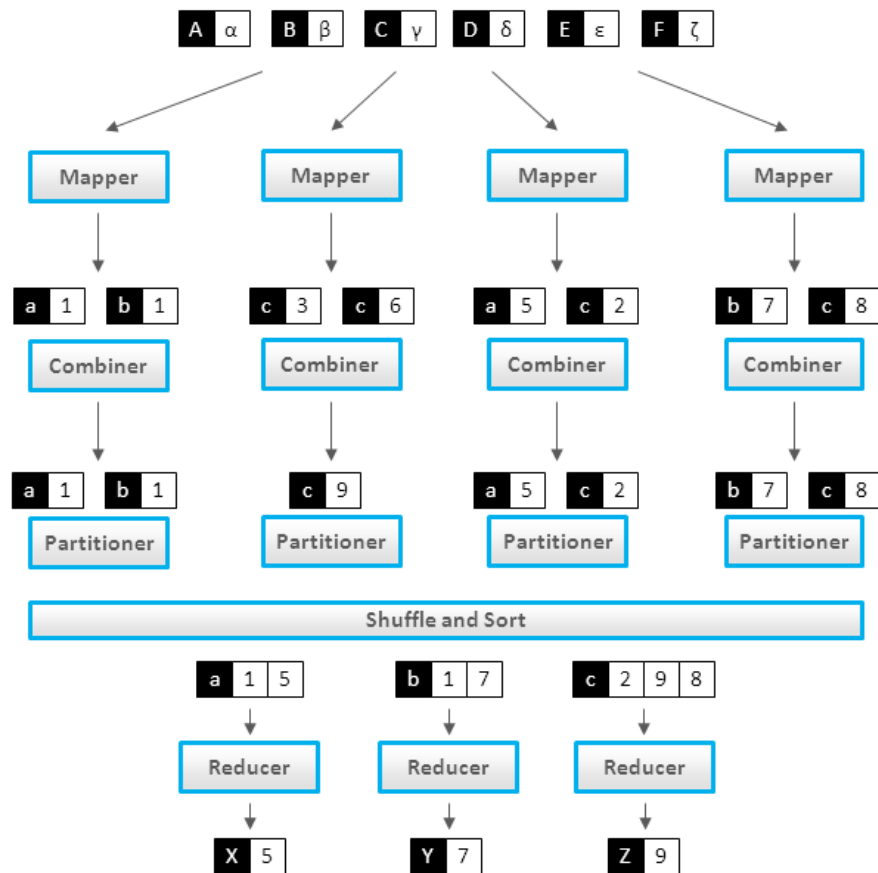
MAPREDUCE

MapReduce involves splitting a problem into subtasks and processing these subtasks in parallel.

This takes place in (approximately) two phases:

1. the mapper phase
 - shuffle/sort
2. the reducer phase

MAP-REDUCE



MAPREDUCE

Additional intermediate steps in a MapReduce workflow:

- mappers – filter & transform data
- combiners – perform reducer operations on the mapper node (optional step, to reduce network traffic and disk I/O).
- partitioners – shuffle/sort/redirect mapper output
- reducers – aggregate results

MAPREDUCE

The MapReduce framework handles a lot of messy details for you:

- parallelization & distribution (eg, input splitting)
- partitioning (shuffle/sort/redirect)
- fault-tolerance (fact: tasks/nodes will fail!)
- I/O scheduling
- status and monitoring

This (along with the functional semantics) allows you to focus on solving the problem instead of accounting & housekeeping details.

MAPREDUCE: HADOOP

Hadoop is a popular open-source Java-based implementation of the map-reduce framework (including file storage for input/output).

You can download Hadoop and configure a set of machines to operate as a MapReduce cluster, or you can run it as a service via Amazon's Elastic MapReduce.

Hadoop is written in Java, but the Hadoop Streaming utility allows client code to be supplied as executables (eg, written in any language)

MAPREDUCE: HADOOP

Hadoop Distributed File System (HDFS):

- Data is replicated in the (distributed) file system across several nodes
- This permits locality optimization (and fault tolerance) by allowing the mapper tasks to run on the same nodes where the data resides
- So we move code to data (instead of data to code), thus avoiding a lot of network traffic and disk I/O

MAPREDUCE

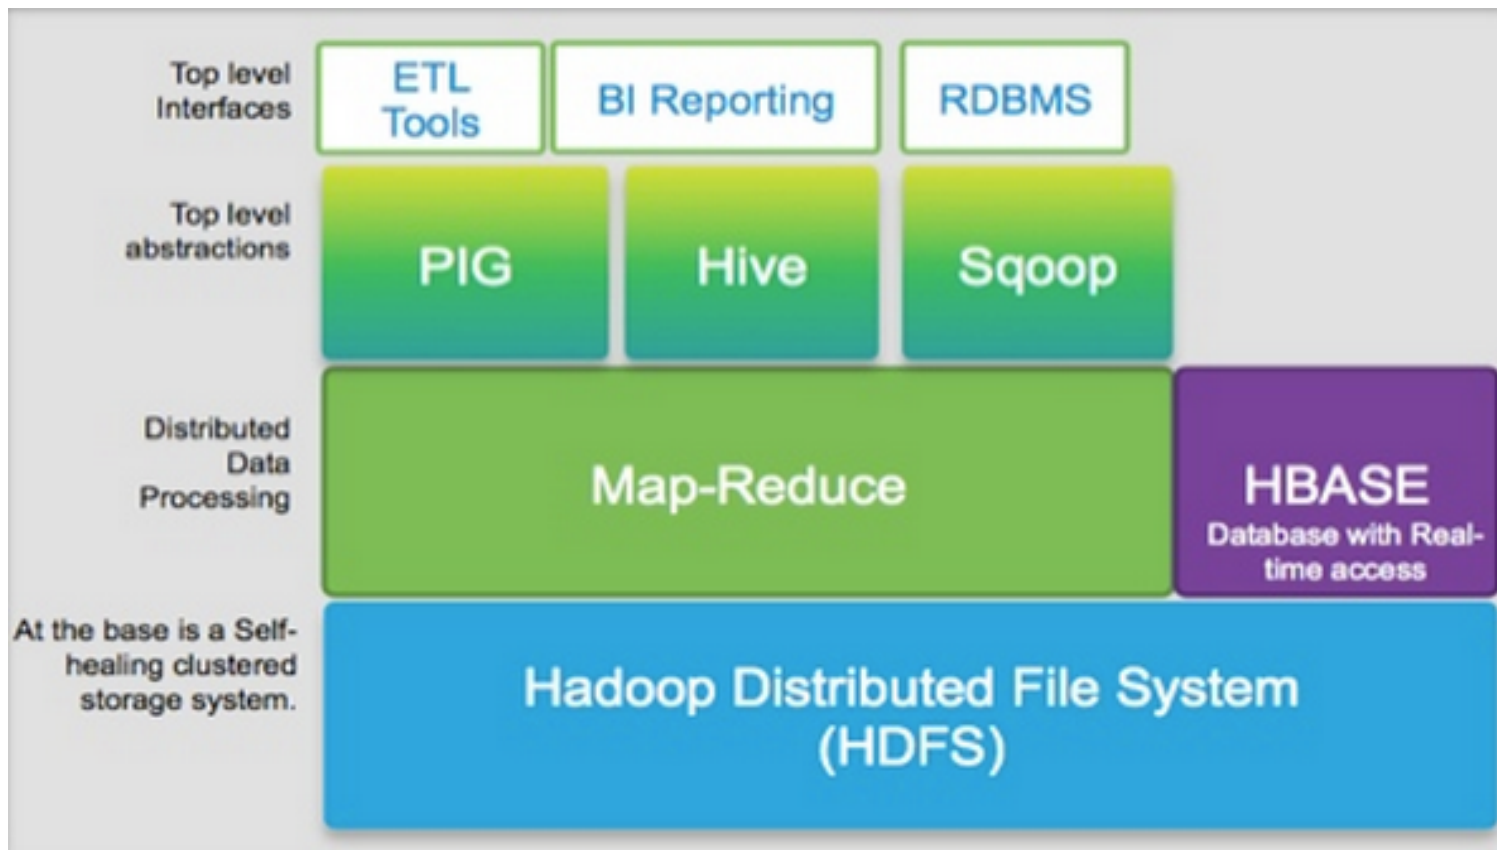
It's possible to overlay the map-reduce framework with an additional declarative syntax.

This makes operations like `SELECT` & `JOIN` easier to implement and less error prone.

Popular examples include ***Pig*** and ***Hive***.



MAPREDUCE: HADOOP ECOSYSTEM

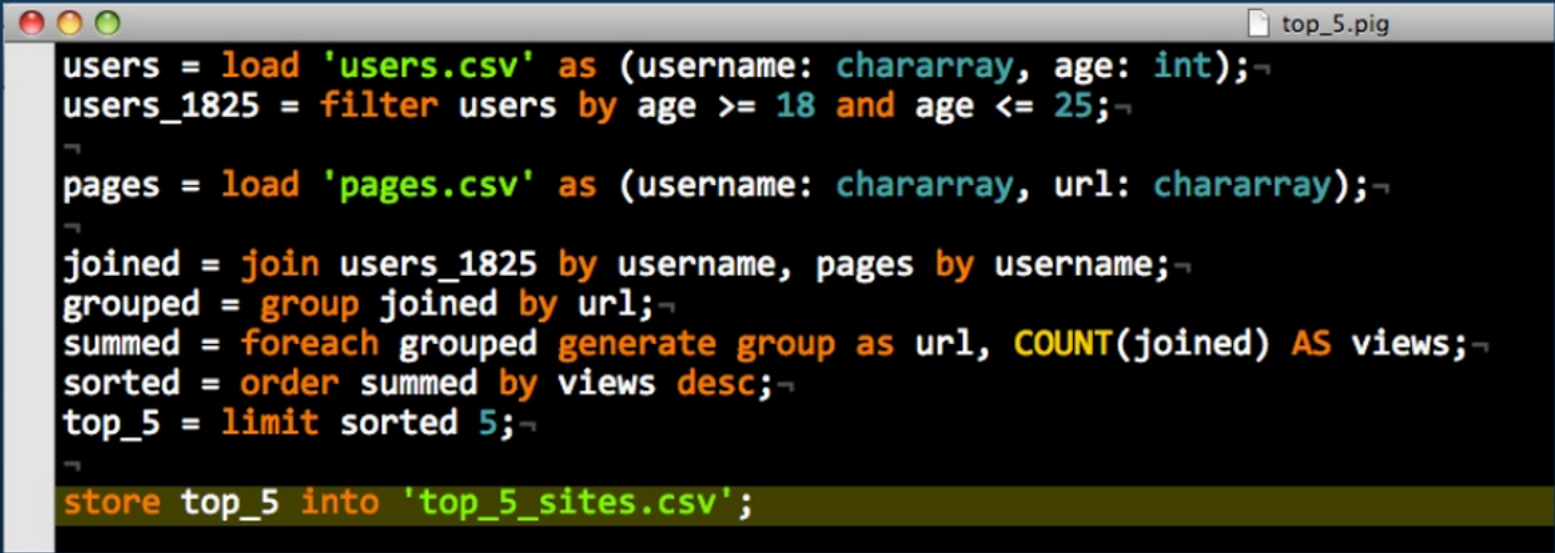


Why Pig?

- ▶ Because I bet you can read the following script.

MAPREDUCE

A Real Pig Script



```
users = load 'users.csv' as (username: chararray, age: int);
users_1825 = filter users by age >= 18 and age <= 25;

pages = load 'pages.csv' as (username: chararray, url: chararray);

joined = join users_1825 by username, pages by username;
grouped = group joined by url;
summed = foreach grouped generate group as url, COUNT(joined) AS views;
sorted = order summed by views desc;
top_5 = limit sorted 5;

store top_5 into 'top_5_sites.csv';
```

- ▶ Now, just for fun... the same calculation in vanilla Hadoop MapReduce.

MAPREDUCE

No, seriously.

[illegible]

INTRO TO DATA SCIENCE

MAPREDUCE: WORD COUNT

MAPREDUCE: WORD COUNT

MapReduce processes data in terms of key-value pairs:

input $\langle k1, v1 \rangle$

mapper $\langle k1, v1 \rangle \rightarrow \langle k2, v2 \rangle$

(partitioner) $\langle k2, v2 \rangle \rightarrow \langle k2, [all\ k2\ values] \rangle$

reducer $\langle k2, [all\ k2\ values] \rangle \rightarrow \langle k3, v3 \rangle$

MAPREDUCE: WORD COUNT

Using the following input, we can implement the “Hello World” of MapReduce: a word count

```
where
where in
where in the
where in the world
where in the world is
where in the world is carmen
where in the world is carmen sandiego
```

MAPREDUCE: WORD COUNT

```
def mapper(doc_id, doc):  
    terms = tokenize(doc)  
    for term in terms:  
        count = 1  
        yield term, count
```

The mapper emits key-value pairs for each word encountered in the input data:

```
where 1  
where 1  
in 1  
where 1  
in 1  
the 1  
...  
<term> 1
```

MAPREDUCE: WORD COUNT

The *partitioner*

- internal to the MapReduce framework, so we don't have to write this ourselves
- shuffles & sorts the mapper output, and redirects all intermediate results for a given key (e.g. term) to a single reducer

| | |
|----------|-----------------------|
| where | [1, 1, 1, 1, 1, 1, 1] |
| in | [1, 1, 1, 1, 1, 1] |
| the | [1, 1, 1, 1, 1] |
| world | [1, 1, 1, 1] |
| is | [1, 1, 1] |
| carmen | [1, 1] |
| sandiego | [1] |

MAPREDUCE: WORD COUNT

```
def reducer(term, counts):  
    total = 0  
    for count in counts:  
        total += 1  
    yield term, total
```

The reducer receives all values for a given key and aggregates (in this case, sums) the results:

| | |
|----------|---|
| carmen | 2 |
| is | 3 |
| in | 6 |
| the | 5 |
| sandiego | 1 |
| where | 7 |
| world | 4 |

INTRO TO DATA SCIENCE

FROM SQL TO HIVE QL

SQL

SQL is a query language designed to load, retrieve, and update data from relational databases

What are some of the most commonly known relational databases?



SQL

SQL is a query language to load, retrieve, and update data in relational databases

What are some of the most commonly known relational databases?

- MySQL
- PostgreSQL
- MS-SQL
- Oracle
- SQLite

SQL

SELECT: Allows you to retrieve information from a table

Syntax:

```
SELECT col1, col2  
FROM table WHERE [some condition]
```

Example:

```
SELECT poll_title, poll_date  
FROM polls  
WHERE romney_pct > obama_pct
```

SQL

GROUP BY: Allows you to aggregate information.

Syntax:

```
SELECT col1, AVG(col2)
FROM table GROUP BY col1
```

Example:


```
SELECT poll_date, AVG(obama_pct)
FROM polls GROUP BY poll_date
```

SQL

JOIN: Couple different kinds in Hive:

JOIN:

produces only the records that match both tables




SQL

JOIN: Couple different kinds in Hive:

FULL OUTER JOIN:

produces all records in both tables, with matching records from both sides where available. If no match, the missing side will contain null




SQL

JOIN: Couple different kinds in Hive:

LEFT OUTER JOIN (RIGHT OUTER JOIN):

produces records from Table A, with matching records (where available) in Table B. If there is no match, the right side will contain null.



SQL

JOIN: Couple different kinds in Hive:

LEFT SEMI JOIN:

produces only records from the left table

equivalent in vanilla SQL:

```
SELECT name FROM table_1 a WHERE EXISTS( SELECT *  
FROM table_2 b WHERE (a.name=b.name))
```

INTRO TO DATA SCIENCE

AWS

AWS

Overview of Services



Compute



Storage



Database



Messaging



Networking



Content
Delivery



Deployment &
Management



Monitoring

AWS

Regions

- us-east-1 (North Virginia)
- us-west-1 (Northern California)
- ap-southeast-1 (Singapore)
- eu-west-1 (Ireland)
- also other regions in Asia, Europe, North and South America

Availability Zones (AZ)

- Each region has several availability zones

AWS

Common Services (for data science):

- S3: Simple Storage Service
- EC2: Elastic Cloud Compute
- EMR: Elastic MapReduce

AWS: S3 (SIMPLE STORAGE SERVICE)

- Sort of a file system
- Buckets ~ Folders, Keys ~ Files
 - Subdirectories aren't really folders, but rather just prefixes to the key name
- Region Specific -- but it usually doesn't matter (except for latency)
- Relatively Cheap
- Complex Permission Structure (about as complex as a real file system)

AWS: EC2 (ELASTIC CLOUD COMPUTE)

Instance Types:

- General Purpose
- Compute Optimized
- Memory Optimized
- Storage Optimized
- GPU Instances

AWS: EC2 (ELASTIC CLOUD COMPUTE)

Example Pricing for On-Demand Linux instances in US East

| | |
|---|--|
| t2.micro <ul style="list-style-type: none">● 1 CPU● 1 GB RAM● \$0.013 per Hour | i2.8xlarge <ul style="list-style-type: none">● 32 CPUs● 244 GB RAM● \$6.82 per Hour |
| m3.xlarge <ul style="list-style-type: none">● 4 CPUs● 15 GB RAM● \$0.28 per Hour | r3.large <ul style="list-style-type: none">● 2 CPUs● 15 GB RAM● \$0.175 per Hour |

Not all instance types available in all regions (e.g. High Storage Instances not available in us-west-1)

AWS: AMI (AMAZON MACHINE IMAGES)

- An **AMI** is a copyable snapshot of an instance's contents and configuration
- If you've ever run a virtual machine (e.g. VirtualBox, VMware, Parallels), you may be familiar
- EC2 instances start as AMIs that are then instantiated
- StarCluster uses AMIs extensively to preconfigure EC2 instances to create a cluster
- A well-configured **AMI** can save a lot of time
- **AMIs** are region specific. Copies to other regions receive a new ID

AWS: EMR (ELASTIC MAPREDUCE)

“Hadoop in the Cloud”

- Uses EC2 instances
- Web interface for provisioning cluster
- Installs Hadoop, etc. automatically
- Installs Hive, Pig, HBase automatically if desired
- Automatically configures Hadoop to be able to read your S3 buckets as if they were part of HDFS
- Charges a small premium on top of EC2 prices